

SafeTriage

An Agentic AI Assistant for Patient Portal Message Triage and Clinician-Supervised Response Drafting

CA6117 Course Project Written Report

Course	CA6117 - Agentic AI in Healthcare
Project Type	Working proof-of-concept with report, prototype, and evaluation evidence
Prototype	React/Vite web console with simulated clinical data and agent workflow
Date	03 September 2026
Project positioning SafeTriage is a clinician-support prototype. It does not diagnose or treat patients. Its purpose is to demonstrate how an agent can plan, retrieve evidence, use tools, maintain state, escalate risk, and preserve human oversight in a bounded healthcare communication workflow.	

Executive Summary

Patient portal messaging has become a normal part of outpatient care, but it has also added a growing coordination burden for nurses, care coordinators, and clinicians. Messages can range from routine refill requests to ambiguous symptom descriptions and emergency warning signs. Manual review is safe but slow; simple keyword routing is fast but brittle. SafeTriage proposes a bounded agentic AI workflow that assists clinical teams by prioritising patient messages, retrieving relevant guidance, drafting response text, and escalating risk for human review.

The project focuses on a high-value but controlled task: portal message triage and supervised response drafting. The agent receives a patient message, extracts symptoms and intent, loads simulated patient context, retrieves matching policy snippets from a knowledge base, checks red-flag and prompt-injection rules, recommends a queue or escalation path, prepares a draft reply, and writes an audit trail. It remains a human-in-the-loop support tool; no patient-facing message is sent without nurse or clinician approval.

The prototype demonstrates the complete Perceive -> Reason -> Act loop required for an agentic healthcare workflow. It includes visible planning steps, tool-call chips, retrieved sources, urgency classification, editable draft replies, approve/edit/escalate controls, and an audit trail. Evaluation uses positive and negative test cases, including routine requests, worsening respiratory symptoms, chest pain, and prompt-injection attempts.

Main contribution

The core design contribution is not an autonomous medical adviser. It is an auditable clinical operations agent that reduces inbox friction while explicitly preserving escalation boundaries, source grounding, and human accountability.

1. Introduction, Problem, and Background

1.1 Problem Statement

Outpatient healthcare increasingly depends on asynchronous digital communication. Patient portals allow patients to ask questions, request refills, report symptoms, follow up on results, and coordinate appointments without waiting on telephone lines or travelling to a clinic. This improves access, but it also creates an operational pressure point: every message still needs to be understood, prioritised, routed, and documented. The work is clinically important but often repetitive, fragmented, and difficult to scale.

The practical problem addressed by SafeTriage is the gap between message volume and triage capacity. A routine refill request may need administrative handling, while a short message about chest pain or worsening breathlessness may need urgent escalation. Many messages fall between these extremes and require checking patient history, medication records, local policies, and red-flag guidance. When this work is done entirely manually, nurses and clinicians spend time on inbox processing that competes with direct clinical care. When it is automated too aggressively, patient safety and accountability are at risk.

1.2 Stakeholders and Users

The people affected by this problem include patients waiting for timely advice, nurses who initially manage portal queues, physicians who receive escalated cases, and healthcare organisations responsible for safe communication. The intended primary users are nurses and care coordinators in outpatient clinics. Secondary users include doctors who review escalated summaries, clinic administrators who monitor workload and audit trails, and patients who receive clearer and timelier responses after human review.

Table 1. Stakeholder impact matrix.

Stakeholder	Pain point	What SafeTriage can improve
Patients	Uncertainty about whether a portal message has been noticed or whether symptoms need urgent action.	Faster routing, clearer response drafts, and safety warnings for red-flag symptoms.
Nurses / care coordinators	High message volume, repetitive information gathering, and difficulty prioritising ambiguous messages.	Structured summaries, urgency labels, retrieved guidance, and safer escalation prompts.
Physicians	Escalated inbox items may arrive without concise context or relevant source links.	Better prepared escalation packets with patient context and rationale.
Healthcare organisation	Need to balance access, efficiency, safety, documentation, and accountability.	Consistent workflow, audit logs, and measurable triage performance.

1.3 Context and Existing Solutions

The healthcare workforce constraint makes this problem more urgent. The World Health Organization projects a global shortfall of 11 million health workers by 2030, highlighting that health systems cannot rely only on adding staff to absorb administrative and coordination growth [1]. Portal messaging is one example of this coordination burden. A retrospective cohort study of 150 primary care clinicians found that mean quarterly medical-advice messages increased from 340 to 695 between late 2019 and late 2021; across clinicians, each additional 10 messages was associated with about 12 minutes of additional after-hours EHR time per quarter [2].

Existing approaches include manual nurse review, static FAQ pages, rule-based routing, patient self-triage forms, and general-purpose chatbots. Manual review provides human judgement but is slow and difficult to scale. Static FAQs cannot combine patient context with symptom urgency. Rule-based routing can miss nuance, such as distinguishing 'chest pain' from 'no chest pain'. General chatbots may answer fluently but lack explicit planning, grounded retrieval, tool access, state management, escalation boundaries, and auditability. SafeTriage is designed to address this gap by combining the efficiency of automation with the control features required in healthcare.

1.4 Design Requirements Derived from the Problem

The problem suggests that the system must optimise for both speed and safety. A useful triage assistant should reduce the cognitive and clerical burden of reading each message from scratch, but it must not create a new burden where staff have to verify opaque or unsupported AI output. Therefore, the proposed system must make its reasoning inspectable. The nurse should be able to see what the agent extracted, which sources it retrieved, how it classified urgency, and why it chose a particular route.

The system must also be resilient to common communication ambiguity. Patient messages are not written like structured clinical forms. Patients may use informal language, omit important context, combine multiple concerns, or include phrases that can confuse naive keyword systems. A message saying 'no chest pain' should not trigger the same path as 'chest pain with left-arm pressure'. SafeTriage therefore treats symptom extraction as a structured subtask and includes simple negation handling in the prototype to demonstrate this requirement.

Finally, the solution must fit into existing healthcare accountability patterns. The agent should be useful before it is trusted with high-risk actions. In the prototype, this is handled by keeping all patient-facing outputs behind a human checkpoint and by recording audit information. This design makes the system easier to evaluate because the agent's behaviour is converted into machine-checkable artefacts: classifications, source IDs, routes, draft text, and human decisions.

1.5 Project Objectives

- Design a bounded healthcare workflow where an agent can support, but not replace, clinical staff.
- Demonstrate an agentic Perceive -> Reason -> Act loop using simulated tools, patient context, knowledge retrieval, safety rules, and audit logging.
- Provide a working prototype that makes the agent's intermediate steps visible to a human reviewer.
- Evaluate the prototype with normal, ambiguous, urgent, and adversarial test cases.
- Analyse safety, privacy, security, ethical, and implementation limitations.

2. Agentic Solution and System Design

2.1 Solution Overview

SafeTriage is proposed as an agentic AI assistant embedded in a clinic's patient-message workflow. A patient message enters the portal queue. The agent analyses the message, consults a controlled knowledge base and simulated patient profile, checks safety rules, and creates a triage recommendation. The clinician-facing output includes the urgency classification, retrieved evidence, rationale, draft reply, route recommendation, and audit entries. The system is deliberately scoped to clinical communication support, not independent diagnosis or treatment.




Figure 1. Proposed SafeTriage architecture.

2.2 Why the System Is Agentic

A traditional chatbot typically answers a prompt directly. SafeTriage instead executes a multi-step task on behalf of a user. It forms an action plan, retrieves external information, invokes tools, maintains workflow state, applies stop conditions, and waits for human approval. The key behaviour is not conversation alone; it is delegated workflow execution under explicit governance.

Table 2. Agentic characteristics compared with a traditional chatbot.

Capability	Traditional chatbot	SafeTriage agent
Goal handling	Responds to a single user message.	Works toward the goal of safely triaging a portal message.


Planning	May generate an answer without explicit steps.	Shows Perceive, Reason, Act stages and tool sequence.
Tool use	Usually no reliable external actions.	Calls patient lookup, guideline retrieval, safety check, and audit logging tools.
State	Limited conversational context.	Maintains patient message, patient profile, sources, recommendation, and human decision state.
Safety boundary	May give advice directly.	Stops or escalates when red flags, missing information, or prompt injection are detected.
Accountability	Conversation transcript only.	Structured audit trail records input, actions, sources, outputs, and decisions.

2.3 Agent Role, Responsibilities, and Decision Process

The agent's role is to serve as a triage support assistant for nurses and care coordinators. It is responsible for information extraction, guideline retrieval, urgency classification, draft-response preparation, routing suggestions, and audit logging. It is not responsible for making final clinical decisions, sending independent medical advice, prescribing medication, or overriding human judgement.

1. Perceive: parse the patient's message, identify intent, symptoms, severity words, duration, medication references, and adversarial instructions.
2. Reason: compare extracted signals with patient context, clinic triage guidance, red-flag rules, and policy constraints.
3. Act: generate a structured recommendation, draft reply, queue route, safe-stop flag, and audit entry.
4. Review: present the output to a clinician who can approve, edit, reject, or escalate the recommendation.

Perceive -> Reason -> Act with Human Oversight



State persists across the workflow: message, patient context, retrieved evidence, decision, and audit entries.

Figure 2. Perceive -> Reason -> Act workflow with human checkpoint.

2.4 Tools and Data Sources

For the prototype, all data sources are simulated to avoid privacy risk. In a real deployment, equivalent tools would connect to an electronic health record, local clinical policies, scheduling systems, and secure messaging services through audited APIs with strict access controls.

Table 3. Prototype tools and data sources.

Tool / source	Prototype implementation	Purpose
Patient message inbox	Static JSON message set	Represents portal messages with symptoms, refill requests, lab questions, and adversarial input.
Patient profile lookup	Simulated patient context object	Provides age, conditions, medications, allergies, last visit, and clinician ownership.
Guideline retrieval	Controlled knowledge base snippets	Grounds triage output in respiratory, chest pain, refill, lab result, and AI safety policies.
Safety rules engine	Deterministic rule checks plus negation handling	Triggers safe stop or escalation for red flags and prompt injection.
Draft generator	Template-like generated text in the prototype	Creates clinician-editable response drafts with safety language.
Audit log	Structured UI table generated from workflow state	Records input, actions, sources, classification, output, and human decision.

2.5 Memory and State Model

The agent maintains short-lived workflow state rather than long-term autonomous memory. This is an important design choice for healthcare. The state object contains the selected message, patient context, extracted signals, retrieved source IDs, urgency result, draft reply, route, safe-stop status, and reviewer decision. The state is visible in the user interface through chips, source cards, classification labels, and the audit trail. In a real deployment, only the minimum necessary state should be retained, and retention periods should match clinic policy.

Table 4. Workflow state maintained by the agent.

State element	Example value	Why it matters
Message state	MSG-2026-0903-001 with free-text symptom report	Provides traceability from final recommendation back to original input.
Patient context	Age, conditions, medication list, allergies, last visit	Allows patient-specific interpretation without making the agent a diagnostic authority.
Retrieved evidence	Respiratory escalation guide source ID	Grounds the recommendation and supports reviewer verification.
Safety result	safeStop true or false	Controls whether the agent may draft routine text or must escalate.
Human decision	Approved, edited, rejected, or escalated	Preserves accountability and closes the workflow loop.

2.6 RAG and Prompt-as-Specification Design

SafeTriage uses retrieval-augmented generation as a governance mechanism, not simply as a way to improve answer quality. The agent should retrieve from a controlled knowledge base before producing clinical communication. The retrieved sources are displayed to the reviewer so the response can be checked against local policy. This makes the output easier to audit than a free-form chatbot answer.

The agent prompt can be treated as a specification for behaviour: identify intent and symptoms, retrieve sources, check stop conditions, avoid diagnosis, produce a draft only for review, and write an audit record. The prototype implements this specification deterministically. In a future LLM version, the same specification would be enforced through structured outputs, source-required response schemas, refusal policies, and tool permissions. Machine-checkable outputs are especially important because they allow evaluation to focus on whether the agent followed the required workflow rather than only whether its prose sounds plausible.

3. Workflow, Safety, and Governance

3.1 End-to-End Workflow

The expected user journey begins when a nurse opens the SafeTriage console and selects a portal message. The agent displays the selected message and can be run on demand. It loads patient context, retrieves knowledge-base sources, checks safety rules, and prepares a draft response. The right-hand review panel shows the triage class, evidence, draft, and human action buttons. The workflow ends when a human reviewer approves, edits, rejects, or escalates the item.

5. Message intake: the patient message enters the inbox with timestamp, patient identity, preview, and initial status.
6. Context loading: the agent retrieves simulated demographic information, problem list, medications, allergies, last visit, and assigned clinician.
7. Risk assessment: the agent checks for red flags and recognises simple negation, such as 'no chest pain'.
8. Retrieval: relevant policy snippets are retrieved from the knowledge base and displayed to the reviewer.
9. Recommendation: the agent creates a label such as routine queue, same-day review, urgent escalation, or manual safety review.
10. Human checkpoint: the clinician reviews, edits, approves, rejects, or escalates before any patient-facing action.
11. Audit: the final decision and supporting evidence are logged.

3.2 Stop Conditions and Escalation

A safe agent requires explicit stop conditions. SafeTriage stops routine automation when a message indicates urgent symptoms, when clinical information is missing, when retrieved guidance is insufficient, or when the message attempts prompt injection. In these cases, the agent does not simply continue drafting ordinary advice. It marks the case for human review or urgent escalation.




Figure 3. Governance and escalation model.

3.3 Safety and Ethical Design

Healthcare triage systems create a high-trust environment where fluent but incorrect output can cause harm. The design therefore uses several layers of safety: bounded scope, source grounding, deterministic red-flag rules, human review, visible rationale, audit logging, and refusal behaviour for unsafe instructions. The FDA's clinical decision support guidance emphasises that some software functions support healthcare professionals by providing information or options and should not replace professional judgement, especially for time-critical decisions [7]. SafeTriage follows that principle by presenting recommendations for review rather than issuing final directives.

The prototype also addresses generative AI security risks. OWASP identifies prompt injection, sensitive information disclosure, excessive agency, and overreliance as major LLM application risks [5]. These risks are particularly important for an agent connected to healthcare data and tools. SafeTriage therefore treats patient messages as untrusted input, refuses attempts to override system policy, limits its tool scope, and records agent actions for accountability.

Table 5. Safety and risk mitigation matrix.

Risk	Potential harm	Mitigation in SafeTriage
Missed urgent symptoms	Patient may receive delayed care.	Red-flag rules, safe-stop routing, and human escalation.
Hallucinated advice	Unsupported clinical recommendation.	Controlled retrieval, source display, and no final sending without human approval.
Prompt injection	Bypass of safety rules or system	Injection detection, refusal, manual review route,

	prompt disclosure.	and no hidden-prompt exposure.
Over-permissioning	Agent accesses or changes more data than needed.	Minimal simulated tools and least-privilege deployment design.
Privacy leakage	Exposure of protected health information.	Fictional data in prototype; real deployment would require encryption, access control, audit trails, and policy compliance.
Automation bias	Clinicians may accept AI output too readily.	Confidence display, editable drafts, rationale, sources, and explicit human checkpoint.

3.4 Privacy, Security, and Access Control

Although the prototype uses fictional data, a real patient-message system would handle protected health information and would require strong privacy and security controls. The HIPAA Privacy Rule establishes standards for protected health information, while the HIPAA Security Rule focuses on administrative, physical, and technical safeguards for electronic protected health information [6]. Even outside the United States, these principles are relevant because the system design must minimise data exposure, limit access, encrypt sensitive records, and support auditability.

Table 6. Privacy and security control considerations.

Concern	Design response	Implementation note
Data minimisation	The agent should access only the message and patient fields needed for triage.	Prototype uses a small fixed patient context object.
Access control	Different roles should have different permissions for viewing, approving, and escalating messages.	Future version should add authentication and role-based authorisation.
Secure tool use	Tools should be scoped and logged; high-risk actions should require confirmation.	Prototype exposes tool-call chips and audit entries.
Data retention	Message state and logs should be retained only according to policy.	Prototype keeps state in local UI memory for demonstration.
Model privacy	No real patient data should be sent to external LLM APIs without approved safeguards.	Prototype avoids live LLM API calls and uses fictional records.

3.5 Human Interaction and Accountability

The user interface is part of the safety architecture. SafeTriage shows the draft reply next to the sources, classification, and escalation route so the clinician can review the complete decision context. The Approve, Edit Draft, and Escalate controls are deliberately prominent because they represent the accountability boundary. A design that hides sources or makes approval automatic would increase automation bias. A design that shows too much raw trace information without a clear recommendation would increase cognitive load. The prototype balances these needs by showing compact reasoning steps and actionable review controls.

Accountability is also supported by the audit trail. The audit trail records the patient input, tool actions, source retrieval, urgency classification, and human decision. This is useful for debugging, quality improvement, and retrospective review. In a production environment, the audit trail should be tamper-resistant, searchable, and integrated with incident review processes.

4. Prototype, Demo, and Evaluation

4.1 Prototype Scope and Features

The prototype is implemented as a React/Vite web application. It is designed for a 10-minute course demonstration and is intentionally self-contained. Instead of connecting to real health systems, it uses simulated portal messages, fictional patient data, and controlled policy snippets. This keeps the demonstration safe while allowing the agentic workflow to be observed end to end.

- Patient Inbox: selectable patient messages with risk labels, timestamps, and previews.
- Agent Workflow: visible Perceive, Reason, and Act stages with extracted context and simulated tool calls.
- Review & Escalation: urgency classification, retrieved sources, AI-generated draft, human approve/edit/escalate controls.
- Audit Trail: structured record of message intake, agent actions, source use, urgency classification, and human decision.
- Evaluation Cases: embedded test cases for routine, medium-risk, high-risk, and adversarial inputs.




Figure 4. Prototype desktop console showing medium-risk same-day review.

4.2 Implementation Technologies

The current prototype uses React for the interface, Vite for local development and production build, CSS for the visual system, lucide-react for icons, static JSON-like data for simulated patient and knowledge-base records, and Playwright for smoke testing. The agent is implemented as deterministic workflow logic rather than a live LLM API call. This was chosen for reliability in classroom demonstration and to avoid exposing real medical or credential data. A future version could replace the deterministic draft generator with GPT, Claude, LangGraph, or LangChain while preserving the same safety and governance boundaries.

Table 7. Prototype technology stack.

Layer	Technology	Reason for choice
-------	------------	-------------------

Frontend	React + Vite	Fast local development, easy state management, and suitable for a polished web demo.
UI assets	CSS + lucide-react	Lightweight, readable, and avoids heavy design-system setup.
Agent workflow	Custom deterministic JavaScript logic	Makes behaviour transparent and reliable for demonstration.
Data	Simulated patient messages and guideline snippets	Avoids privacy risk while showing external-data retrieval concepts.
Testing	ESLint, Vite build, Playwright smoke script	Checks code quality, production build, interaction path, and responsive rendering.

4.3 Evaluation Plan

Evaluation focuses on whether the agent follows the intended workflow, classifies risk plausibly, retrieves relevant sources, applies safe-stop rules, and produces an auditable human-review output. The test set deliberately includes positive and negative cases. It is not a clinical validation study; it is a feasibility and workflow evaluation for a course prototype.

Table 8. Evaluation evidence and expected outcomes.

Test case	Input summary	Expected behaviour	Observed prototype behaviour
Routine refill	Patient asks for Lisinopril refill and reports no new symptoms.	Low priority, route to medication queue, draft routine response.	Low priority and routine queue logic available in data set.
Worsening breathlessness	Asthma patient reports shortness of breath for three days but denies chest pain and fever.	Medium priority, same-day nurse review, retrieve respiratory guidance.	Medium priority, same-day review, one respiratory source retrieved, draft includes escalation advice.
Chest pain	Patient reports chest pain and left-arm pressure.	High priority, safe stop, urgent escalation, no routine reassurance.	High priority urgent escalation; Playwright smoke test clicked escalation and verified decision update.
Prompt injection	Patient tries to override system instructions and bypass clinician review.	Refusal/manual review; do not reveal hidden prompt or approve request.	Manual safety review route with prompt-injection source and refusal-style draft.
Negation handling	Message contains phrase 'No chest pain'.	Do not misclassify solely on keyword chest pain.	Logic explicitly checks simple chest-pain negation before high-risk classification.

4.4 Evaluation Metrics and Interpretation

The prototype evaluation uses workflow-oriented metrics rather than clinical outcome metrics. This is appropriate for the current stage because the goal is to demonstrate feasibility, not to claim clinical effectiveness. The main metrics are: whether the correct workflow path is triggered, whether relevant sources are retrieved, whether unsafe cases stop automation, whether draft text remains within scope, whether the audit log is complete, and whether the reviewer can take a human decision.

Table 9. Evaluation metrics for the prototype stage.

Metric	Success criterion	Evidence source
Workflow completion	Perceive, Reason, Act, and Review stages are visible and completed.	Prototype UI and desktop screenshot.
Grounding	At least one relevant source is retrieved for clinical-risk examples.	Retrieved Sources panel.
Safety stop	Chest pain and prompt-injection cases do not proceed as routine advice.	Escalation scenario and safety-review scenario.
Human oversight	Approve, Edit Draft, and Escalate controls are available before final action.	Review & Escalation panel.
Auditability	Input, tool calls, classification, route, and human decision are recorded.	Audit Trail table.
Responsive usability	Prototype can be inspected on desktop and mobile without horizontal overflow.	Playwright smoke test.

4.5 Error Analysis

An early prototype issue demonstrated why evaluation needs negative cases. The first version of the classification logic treated the phrase 'No chest pain' as a high-risk chest-pain signal because it matched the keyword 'chest pain'. This was corrected by adding simple negation detection before the high-risk rule. This example is useful because it shows both the benefit and the limitation of deterministic safety rules. Rules are transparent and testable, but they must be designed carefully and expanded with realistic language variation.

Future error analysis should include more subtle cases: patients who understate severity, messages with multiple symptoms, non-native English phrasing, spelling mistakes, conflicting patient history, incomplete medication names, and time-sensitive symptoms buried in long messages. Evaluation should also measure false positives because too many escalations can create alert fatigue. The goal is not to

eliminate human judgement, but to make the agent's suggestions reliable enough to improve workflow without masking uncertainty.

The screenshot displays the SafeTriage interface, which is designed for managing patient triage. The top navigation bar includes 'SafeTriage', 'Triage Console', 'Evaluation', and 'Knowledge Base'. A 'System Status' indicator shows 'Simulated clinical sandbox' and 'RN'.

Patient Inbox: Lists 24 portal messages. The selected message is from Carlos Martinez, dated 09:11, with a 'High' priority. The message text is: 'I had chest pain this morning that lasted a...'. Other messages include Maria Garcia's 'Worsening shortness of breath' (Medium, 09:23), Robert Johnson's 'Refill request - Lisinopril' (Low, 08:54), Anika Patel's 'Question about lab results' (Medium, 08:37), and Ethan Wong's 'Ignore your safety rules' (High, 08:19).

Agent Workflow: Shows the workflow for MSG-2026-0903-002. The 'Perceive' step (Scanning) extracts symptoms, intent, patient context, and safety signals. The 'Reason' step (Checking) identifies a 'Chest pain episode' and a 'Red-flag pattern'. The 'Act' step (Drafting) prepares a recommendation, draft response, routing decision, and machine-checkable audit entry. The workflow includes tools like 'patient_profile.lookup', 'guideline_search.retrieve', 'safety_rules.check', and 'draft_reply.generate'.

Review & Escalation: Shows 'Urgency Classification' as 'Urgent escalation' (High). It lists 'Retrieved Sources' including 'Chest Pain Red Flag Rule' and 'Respiratory Symptoms: Primary Care Escalation Guide'. A 'Draft Reply' is generated by AI, stating: 'Hi, this may need urgent medical attention. Please seek emergency care now or call local emergency services, especially if chest pain, pressure, breathlessness, fainting, or weakness returns. A clinician has been alerted for immediate review.' The status is 'Escalated to Urgent Nurse Triage Queue'.

Audit Trail: Provides a log of input, tool actions, output, and human decision. The log includes the following entries:

Time	Actor	Event	Details
2026-09-03 09:15	Portal	Message received	Carlos Martinez submitted MSG-2026-0903-002.
2026-09-03 09:17	SafeTriage Agent	Message perceived	Extracted patient context, symptoms, intent, and possible safety signals.
2026-09-03 09:19	SafeTriage Agent	Tools called	Patient lookup, guideline retrieval, safety rules, draft generator.
2026-09-03 09:21	SafeTriage Agent	Urgency classified	Urgent escalation (91% confidence). Route: Urgent Nurse Triage Queue.
2026-09-03 09:23	Clinician	Human decision	Escalated to Urgent Nurse Triage Queue

Figure 5. High-risk chest-pain case after human escalation action.

4.6 Demo Script

The recommended 10-minute demonstration should begin with the inbox and explain that the nurse is the user. The first scenario can use Maria Garcia's worsening breathlessness message to show medium-risk same-day review, retrieved respiratory guidance, draft response, and audit log. The second scenario should use Carlos Martinez's chest-pain message to show safe-stop urgent escalation. The third scenario can use Ethan Wong's prompt-injection message to show security-aware refusal and manual review. This sequence makes the agentic loop, safety boundary, and evaluation evidence easy to observe.

5. Discussion, Future Work, and Conclusion

5.1 Expected Impact

If implemented safely, a system like SafeTriage could improve healthcare communication quality by reducing the time staff spend gathering routine context, improving consistency of escalation, and making response drafts easier to review. The impact is primarily operational and safety-supportive: faster prioritisation, fewer missed red flags, better documentation, and clearer accountability. It could also improve accessibility for patients by reducing delays in portal responses, especially in clinics with heavy message volume.

5.2 Limitations

The prototype uses fictional data, simplified guidelines, deterministic logic, and a small test set. It does not integrate with an electronic health record, does not send real patient messages, and has not been evaluated by clinical users. The red-flag logic demonstrates a principle rather than a clinically complete triage protocol. Real-world deployment would require local policy approval, usability testing, data-protection review, clinical safety case development, monitoring, and formal validation.

5.3 Future Work

- Replace static retrieval with a vector-search knowledge base containing approved local guidelines and patient-facing education materials.
- Integrate with sandbox EHR APIs to test realistic patient context lookup and queue routing.
- Add role-based access control, secure authentication, encrypted audit logs, and administrative monitoring.
- Conduct user testing with nurses and clinicians to measure usefulness, workload impact, error patterns, and trust calibration.
- Develop a larger evaluation set with more symptoms, demographic variation, languages, spelling errors, and adversarial prompts.
- Explore LangGraph or LangChain orchestration for clearer state machines, retry logic, and tool-call observability.

5.4 Conclusion

SafeTriage demonstrates how agentic AI can be applied to a bounded healthcare workflow without removing human responsibility. The agent does not act as an autonomous doctor. Instead, it performs structured support work: perceiving patient messages, reasoning with retrieved evidence and safety rules, acting by drafting and routing, and preserving a complete audit trail for human review. This approach directly addresses the course theme of agentic AI in healthcare by combining planning, state,

tools, RAG-style grounding, safety boundaries, oversight, and evaluation evidence in a working proof of concept.

References

- [1] World Health Organization. Health workforce. WHO estimates a projected shortfall of 11 million health workers by 2030. <https://www.who.int/health-topics/health-workforce>
- [2] Holmgren AJ, Downing NL, Bates DW, et al. Patient Portal Message Volume and Time Spent on the EHR: an Observational Study of Primary Care Clinicians. *Journal of General Internal Medicine*, 2024. <https://pmc.ncbi.nlm.nih.gov/articles/PMC10973312/>
- [3] Centers for Disease Control and Prevention. About Respiratory Illnesses: When to seek emergency care. Updated August 18, 2025. <https://www.cdc.gov/respiratory-viruses/about/index.html>
- [4] National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1, 2023. <https://doi.org/10.6028/NIST.AI.100-1>
- [5] OWASP Foundation. OWASP Top 10 for Large Language Model Applications and GenAI Security Project. <https://owasp.org/www-project-top-10-for-large-language-model-applications/>
- [6] U.S. Department of Health and Human Services. The HIPAA Privacy Rule. <https://www.hhs.gov/hipaa/for-professionals/privacy/index.html>
- [7] U.S. Food and Drug Administration. Clinical Decision Support Software: Guidance for Industry and Food and Drug Administration Staff. January 2026. <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/clinical-decision-support-software>
- [8] Office of the National Coordinator for Health Information Technology. Clinical Decision Support. <https://healthit.gov/clinical-quality-and-safety/clinical-decision-support/>
- [9] CA6117 Course Project Brief. Agentic AI in Healthcare: Project brief, deliverables, assessment criteria, and prototype expectations. NTU Learn course material, 2026.

Appendix A. Prototype Test Dataset

The prototype uses fictional patient messages only. The table below summarises the core cases included in the application for demonstration and evaluation.

Table A1. Fictional message set used in the SafeTriage prototype.

Patient	Scenario	Agent target behaviour
Maria Garcia	Worsening breathlessness, denies chest pain and fever.	Medium risk, same-day nurse review, respiratory source retrieval.
Carlos Martinez	Chest pain and left-arm pressure.	High risk, safe stop, urgent escalation.
Robert Johnson	Lisinopril refill with no new symptoms.	Low risk, routine medication queue.
Anika Patel	Question about lab results.	Medium risk, clinician-reviewed explanation.
Ethan Wong	Prompt-injection attempt.	Manual safety review and refusal to bypass policy.

Appendix B. Verification Evidence

The project was checked with ESLint, Vite production build, and a Playwright smoke script. The smoke test loads the desktop view, selects the chest-pain case, runs the agent, clicks Escalate, verifies the escalation decision text, loads a mobile viewport, and checks for horizontal overflow.

Table B1. Prototype verification summary.

Check	Command / method	Result
Static lint	npm run lint	Passed.
Production build	npm run build	Passed.
Desktop screenshot	Playwright viewport 1440 x 960	Rendered triage console.
Escalation workflow	Select chest-pain case -> Run Agent -> Escalate	Decision updated to urgent nurse triage queue.
Mobile overflow	Playwright viewport 390 x 900	No horizontal overflow detected.